

Spatial rescue in biofilms: a double-edged sword

May 16, 2026

Abstract

A biofilm under antibiotic treatment has a clear spatial asymmetry: cells on the surface are exposed to the drug, while cells in the deeper interior are shielded from it. We ask whether that shielded interior makes the population more likely to evolve drug resistance and survive, a process known as evolutionary rescue. Comparing a spatially structured biofilm model against a well-mixed population of the same initial size, and sweeping the antibiotic dose along with the metabolic activity of the interior, we find that the biofilm is a double-edged sword; it rescues less often than the well-mixed reference at low dose, but more often at high dose, and only when the interior is metabolically active. We explain this pattern as a consequence of two opposing effects of the spatial structure: a supply mechanism that generates resistant mutants in the interior yet is subject to drift since it is not exposed to the drug, and a delivery mechanism that transports those mutants to the surface where they can grow and establish but is tied to antibiotic dose via the velocity at which the biofilm deteriorates. Our results suggest that the spatial structure of biofilms can both hinder and facilitate evolutionary rescue, and that the balance between these effects depends on the antibiotic dose and the metabolic activity of the interior.

1 Introduction

Antibiotic resistance is an escalating clinical problem, and biofilms (dense, surface-attached bacterial communities) carry a disproportionate share of it. Biofilms are responsible for a large fraction of chronic infections and are notoriously hard to clear with antibiotics; standard regimens that work in well-mixed culture often fail when the target population is organized as a biofilm. This has driven interest in behavior-modifying interventions, in particular compounds that induce bacterial dispersal away from biofilms toward a planktonic, or more well-mixed state. Such interventions do not kill cells; they simply change how cells are arranged in space. The hope is that disrupting biofilm spatial structure reduces the population's capacity to evolve resistance under antibiotic pressure, essentially disarming the bacteria such that they will be more easily cleared by traditional antibiotic treatments and the host immune system. Knowing whether that hope is justified requires understanding what the spatial structure of a biofilm actually does to a population's evolutionary dynamics in the first place. Does biofilm formation help the population evolve resistance, hurt it, or some combination of both?

Biofilms have a clear spatial asymmetry: cells on the outer surface are exposed to the drug, while cells in the deeper interior are shielded from it. The shielded interior is a spatial refuge, a region of the population where the antibiotic selection pressure does not reach. We ask whether

that refuge makes the population more likely to undergo *evolutionary rescue*, the outcome in which a declining population avoids extinction because a resistant mutant arises and takes over in time. Concretely: does a biofilm with a shielded interior rescue more often than a well-mixed population of the same initial size?

The naive prediction is that the refuge provided by biofilm structures should help facilitate evolutionary rescue. The biofilm creates a gradient of antibiotic concentration, with the shielded interior experiencing lower levels of the drug. This shielded interior lets the population persist for longer under treatment, and a longer overall population-level lifespan means more cell divisions, which means more opportunities for a resistant mutation to arise. More mutations should mean more rescue. But this prediction has at least two distinct ways it could be wrong. First, a mutation that arises in the shielded interior does nothing on its own: there is no selection pressure inside the refuge to make a resistant cell take over, so a core mutation matters only if it eventually reaches the drug-exposed surface in time to grow under selection. The refuge might generate mutations the population can never actually use. Second, the very mechanism that maintains the refuge, namely the replacement of dying surface cells by cells flowing outward from the interior, could act against a new resistant cell trying to establish on the surface by ensuring it is surrounded by a continuously refreshed wild-type majority. The same spatial structure that extends the population's lifetime may also make each individual mutation that arises less useful. Whether the biofilm is a net help or a net hindrance therefore depends on how these effects balance.

We measure that balance directly by simulating a minimum-viable biofilm (a one-dimensional chain of cells with a drug-exposed edge and a shielded core) under antibiotic pressure, and compare its rescue probability against an otherwise identical well-mixed population. By varying the antibiotic dose and the metabolic activity of the interior independently, we ask whether the biofilm's spatial structure is uniformly helpful, uniformly harmful, or, as the results below will show, contingent on both the selection pressure and on what is happening inside the refuge.

2 The model

We model the biofilm as a one-dimensional chain of cells. Real biofilms are three-dimensional, but the central asymmetry we are after (drug-exposed surface versus shielded interior) does not need three dimensions to capture. One dimension keeps the dynamics tractable while preserving the feature that matters: cells at different positions face different antibiotic concentrations.

2.1 The chain and rates

The state of the population at any time t is an ordered list of cells indexed $0, 1, \dots, N(t) - 1$. Index 0 is the deep interior; index $N(t) - 1$ is the outer surface. Each cell has a genotype, either wild-type (WT) or resistant (R), and a position determined by its index. As per convention with evolutionary rescue, the population starts at $t = 0$ with $N(0) = N_0$ WT cells and no resistant cells. The population then evolves under antibiotic pressure, which begins at time $t = 0$, with births and deaths changing the chain's length and composition over time.

Antibiotic exposure is a step function in position. At any given population size $N(t)$, the rightmost l cells of the chain form the edge, the region where antibiotic concentration is lethal to WT. In this model, we assume l to be a fixed integer parameter describing the width of the edge, or the thickness of the drug-exposed layer. All cells with index $< N(t) - l$ form the core,

the shielded interior. The boundary between edge and core sits at position

$$b(t) = \max(0, N(t) - l), \quad (1)$$

so as the population shrinks, the boundary retreats from the surface toward the interior. When $N(t) \leq l$ the boundary collapses to $b = 0$, and the entire population is edge; no core remains. Additionally, the well-mixed equivalent of this model is the special case where $l \geq N_0$, so the edge is as wide as the entire population and there is no core at all. The boundary is always at $b = 0$ in that case, and every cell is drug-exposed from the start. The flat step-function profile is a stylization of the real, smooth antibiotic gradient through a biofilm. Treating the boundary as sharp is the minimum-viable abstraction that captures the asymmetry; so it is reasonable to assume that a smooth gradient would change quantitative numbers but not the qualitative dynamics we are after.

Each cell has a birth rate and a death rate determined by its compartment and its genotype, summarized in Table 1.

	Core ($i < b$)		Edge ($i \geq b$)	
	birth	death	birth	death
Wild-type (WT)	b_c	d_c	b_e	d_e
Resistant (R)	b_c	d_c	b_R	d_R

Table 1: Per-cell birth and death rates by compartment and genotype, where i is the cell index in the chain, and $b(t)$ is the boundary position between core and edge. The eight compartment-genotype-event combinations that drive the Gillespie dynamics are described in §2.2.

The edge WT birth rate $b_e = 1$ sets the time unit, where time is tracked in generations as the inverse of the birth rate $\frac{1}{b_e} = 1$. The edge WT death rate d_e is the antibiotic dose, swept over values $d_e > b_e$ so WT cells in the edge are always net-dying under treatment. Resistant cells in the edge have $b_R = 1$ and $d_R = 0.6$, giving them a positive net growth rate $s_R = b_R - d_R = 0.4$ once they reach the drug-exposed surface; an R lineage that establishes in the edge grows deterministically to dominance. Notably, however, there are different classes of antibiotic, and some drugs are bacteriostatic rather than bactericidal. For a bacteriostatic drug, the edge WT death rate would be $d_e = b_e = 1$, so WT cells in the edge would be net-neutral rather than net-dying. We do not study that case here, but it is straightforward and possibly interesting to implement by setting $d_e = 1$.¹ Additionally, the model allows multiple R lineages to arise and compete, but the assumptions are such that rescue is defined whenever there are at least $r_{\text{thr}} = 10$ R cells in the edge, so the first lineage or combination of lineages to reach that threshold is what determines rescue. As such, we do not examine the dynamics of clonal interference among R lineages, which would be an interesting topic for future work but is not the focus here.

Cells in the core, regardless of genotype, have balanced rates $b_c = d_c$. The core has no net growth on its own, and WT and R are treated identically there because the shielded interior carries no selection. The default core rate is $b_c = 0.2$, a slow but nonzero pace; the dormant-core limit is $b_c = 0$, where the interior produces no births and no deaths at all. Comparing these two limits is one of the central experiments below.²

¹We would expect that if WT cells were stalled in growth but not dying, rescue dynamics would be different.

²To prevent unbounded post-rescue growth, R births in both compartments are multiplied by a Wilson-type

2.2 Dynamics

Events fire one at a time, sampled as a continuous-time Markov chain via the Gillespie algorithm. At each step we compute the total event rate R_{tot} across all eight compartment-genotype-event combinations in Table 1, advance simulated time by an exponential draw with mean $1/R_{\text{tot}}$, and pick which event fires from a categorical distribution with weights proportional to each event's rate.

A birth event picks a parent uniformly at random within its compartment-genotype subgroup. The daughter is inserted adjacent to the parent (at the parent's position or one position to the right, each with probability $1/2$), and the chain grows by one. A death event picks a victim uniformly at random within its subgroup and removes it; the chain contracts by one.

The boundary $b(t)$ is recomputed after every event according to Eq. (1). When the population grows by one and the boundary moves outward, the cell that now sits at the new boundary position has crossed from edge into core; it is reclassified as a core cell. When the population shrinks and the boundary moves inward, the cell that now sits one position outside the new boundary has crossed from core into edge; it is reclassified as an edge cell. These reclassifications are not separate dynamical events: they are bookkeeping consequences of recomputing the boundary against a chain that has just changed length by one.

This bookkeeping is the *conveyor belt* mechanism the rest of the writeup will refer to. As long as the core has cells in it, every edge death pulls an interior cell out to the surface, and every edge birth pushes an outer cell back into the interior. The result is that the edge holds exactly l cells while the core has anything left to supply; the population is shrinking, but the surface layer is being continuously replaced from below. Figure 1(a, b) illustrates one such reclassification on a small didactic chain: a cell in the same physical position is in the core before an edge death and in the edge after, because the boundary has moved over it.

The conveyor belt produces two qualitatively different regimes of decline. While the core still has cells in it ($N(t) > l$), each edge death is replaced immediately by a core cell, so the edge stays at exactly l cells. Only the l edge cells contribute to net loss, at rate $s_e \cdot l$ per unit time, and the total population falls linearly, where $s_e = d_e - b_e$. The duration of this buffered decline is

$$T_1 = \frac{N_0 - l}{s_e \cdot l}, \quad (2)$$

which is long at low dose (slow edge turnover empties the core gradually) and brief at high dose (fast edge turnover empties the core quickly). Once the core is exhausted ($N(t) \leq l$; Figure 1(c)), the edge is the entire remaining population, and there is no reservoir to draw on. The remaining $\leq l$ cells decay exponentially at rate s_e , just as a well-mixed population would. The transition between these two regimes is sharp on the simulation timescale: it happens at the moment the last core cell is converted to edge. After that point, the dynamics on the chain are indistinguishable from those of a one-compartment well-mixed model with $N = l$.

2.3 Mutations, rescue, and initial condition

Mutations are replication-linked: at each WT birth event, with probability μ , the daughter is an R cell instead of WT; otherwise the daughter inherits the parent's genotype. Mutations are

density factor $\max(0, 1 - N/K)$, with $K = N_0$ throughout. WT births stay density-independent. The factor is approximately 1 during the wild-type decline that we study, so it does not affect the dynamics of interest; it matters only once R has taken over.




Figure 1: The conveyor belt as a consequence of recomputing the boundary, illustrated on a small didactic chain with $N = 9$ initial cells and $l = 4$. The rightmost l cells (orange) form the edge; the rest (grey) form the core. (a) Before the event, the cell drawn with a heavy outline sits in the core, immediately inside the boundary $b = N - l$. (b) After an outer edge cell has died, N has dropped to 8 and the recomputed boundary has moved one position inward; the tracked cell, in the same physical position, is now part of the edge. No cell has been shuffled; the bracket has shifted over it. (c) After enough cumulative edge deaths, N reaches l : the boundary collapses to $b = 0$, the entire chain is edge, and the dynamics enter Phase 2. An edge birth is the symmetric operation, with the boundary moving outward by one position and the innermost edge cell reclassified as core.

irreversible (no back-mutation from R to WT), so once a lineage is R, all of its descendants are R. Because mutations ride on births, they appear at specific spatial positions on the chain (the position of the dividing parent); a mutation born deep in the core enters the world at the core's spatial location and must be physically transported to the edge to come under selection.

Every R lineage that ever appears is tagged with a unique identifier when its founder is born, and that identifier propagates to all descendants. This lets us follow each mutation individually: where it was born (core or edge), when it appeared, how its descendant population grew, whether it ever reached the edge, and whether it contributed to rescue. The per-lineage records produced by this tracking are the basis for most of the analyses in later sections.

A run is declared *rescued* when the edge contains at least $r_{\text{thr}} = 10$ resistant cells simultaneously. The choice of $r_{\text{thr}} = 10$ is high enough to be well past the establishment bottleneck (the regime where a small R lineage might still go extinct by chance) but small enough that rescue is declared promptly after takeover begins. Once an R lineage is past establishment in the edge, where $s_R > 0$, it grows deterministically to dominance. Every simulation terminates naturally on one of two events: rescue ($\geq r_{\text{thr}}$ R cells in the edge) or extinction ($N = 0$). There are no artificial time or generation caps. Both endpoints are guaranteed by the dynamics: either a successful resistant lineage establishes in time, or the wild-type population is driven to zero. The

model cannot stall in between.

Every simulation starts at $t = 0$ with $N = N_0$ wild-type cells and no resistant cells. Before $t = 0$ the population is at quasi-steady state, with balanced birth and death rates in both compartments ($b_e = d_e = 1$ pre-treatment in the edge, $b_c = d_c$ in the core), so the biofilm neither grows nor shrinks. At $t = 0$ the antibiotic switches on: the edge WT death rate jumps from $b_e = 1$ to its treatment value $d_e > 1$, and the population begins to decline. What we simulate is therefore the response of an already-mature biofilm to a step-function antibiotic challenge. Real treatments are gradual; this abstraction strips out the loading dynamics so that the rescue dynamics, the only thing we are studying, are not entangled with growth-phase or pharmacokinetic transients.

2.4 The well-mixed reference

To say whether the biofilm helps or hurts, we need a baseline. Our baseline is a well-mixed population of the same initial size N_0 , governed by the edge birth and death rates alone: every cell is drug-exposed, none are shielded. This is the standard setting for evolutionary rescue (Wilson et al., 2017), and it is the natural null model, the same biofilm with all spatial structure removed. We simulate this fully well-mixed model (essentially a one-compartment model and faithful recreation of Wilson’s evolutionary rescue model). Operationally, however, the true well-mixed reference to our biofilm model is the chain model with $l \geq N_0$: when the edge is as wide as the entire population, no cell is ever in the core, the conveyor belt never engages, and the dynamics reduce to a one-compartment birth-death process on scalar counts. We also maintain an independent scalar-count implementation as a cross-check; the two well-mixed models agree statistically within sampling error, which serves as an internal consistency test on the chain model.

2.5 Parameters and simulation details

Unless stated otherwise, all simulations use the parameters in Table 2. Time is measured in edge generations ($b_e = 1$). The dose sweep covers $d_e \in \{1.05, 1.1, 1.2, 1.4, 1.6, 1.9, 2.3, 2.8\}$. The core-activity sweep covers $b_c \in \{0, 0.05, 0.1, 0.2, 0.4, 0.8\}$. The edge-width sweep covers $l \in \{25, 50, 100, 200, 400\}$. Every parameter combination is run with 5,000 independent seeds; the well-mixed reference is run at the same dose points with the same seed count.

Parameter	Value	Meaning
N_0	1000	Initial population size
l	100 (swept)	Fixed value for edge width (# cells exposed to drug)
b_e	1.0	WT birth rate in edge (sets time unit)
d_e	dose (swept)	WT death rate in edge; super-MIC is $d_e > 1$
b_R	1.0	R birth rate in edge
d_R	0.6	R death rate in edge ($s_R = 0.4$)
$b_c = d_c$	0.2 (swept)	Core birth & death rate (0 = dormant)
μ	10^{-4}	Mutation probability per birth
r_{thr}	10	Resistant edge cells required to declare rescue
K	N_0	Density cap for R births (see footnote, §2.1)

Table 2: Default model parameters. The swept variables are dose (d_e), core activity (b_c), and edge width (l).

3 The double-edged sword

Section 1 closed with a naive prediction: the biofilm’s refuge should make rescue more likely than in the well-mixed reference. We also flagged two distinct ways that prediction could be wrong. The geometric refuge could be useless on its own, with the extra lifetime never translating into extra rescues. Or the refuge could be helpful only under specific conditions, perhaps only at high dose or only when the interior is metabolically active, so that an “average” effect across regimes would smear out a sharp contingency. We want to distinguish these failure modes from each other, and from the naive prediction, in a single experiment.

Two knobs, varied independently, do the job. We vary the antibiotic dose between a value just above the minimum inhibitory concentration ($d_e = 1.05$) and a value well above it ($d_e = 2.8$); and we vary the metabolic activity of the interior between $b_c = 0$ (dormant, with no births or deaths inside the refuge) and $b_c = 0.8$ (active, with the interior dividing at a substantial fraction of the edge rate). Each knob isolates a different question. The dormant-core column tests whether the geometric refuge by itself matters: with no births inside, the only thing the biofilm has over the well-mixed reference is extra lifetime, so any increase in rescue probability here would have to come from time alone. The active-core column tests whether what happens inside the refuge matters: the interior is now dividing, adding cell divisions on top of the lifetime extension. Running both at low and high dose then asks whether either effect depends on the selection pressure that is forcing rescue in the first place. In Figure 2 we ran 5,000 independent simulations of the biofilm and 5,000 of the well-mixed reference per parameter combination.

Figure 2 reduces to two observations. First, Panel A shows probability of rescue for both the biofilm and the well-mixed reference at low and high dose, assuming a dormant core ($b_c = 0$), meaning that the interior is not actively dividing or dying (i.e., no turnover of cells in the core). The biofilm rescues less often than the well-mixed reference at both doses, with no hint of a crossover. The same spatial structure that extends the population’s lifetime in a biofilm reduces rather than raises the rescue probability, and it does so at both low and high dose. Thus, extra lifetime alone does not convert into more rescue at any dose we tested.

Second, Panel B shows the same comparison but with an active core ($b_c = 0.8$), meaning that the interior is metabolically active and cells are dividing and dying at a substantial fraction of the edge rate. The dormant-core pattern from Panel A still holds at low dose: the biofilm rescues less often than the well-mixed reference. But at high dose, the pattern flips: the biofilm rescues roughly twice as often as the well-mixed reference. The same spatial structure that reduces the rescue probability at low dose raises it at high dose, allowing biofilms to rescue more often than the well-mixed reference at high dose.

Both observations, especially the second, contradict the naive intuition. The dormant column (Panel A) already shows that extra lifetime alone does not produce more rescue, as the biofilm rescues less often than the well-mixed reference at both low and high dose. The active column (Panel B) then shows that the interior’s metabolic activity is required for the biofilm to rescue more often than the well-mixed reference, and this only happens at high dose (i.e., high d_e), where the antibiotic is most aggressively driving the population down. The same mounting selection pressure that should make rescue harder for everyone is precisely where the biofilm rescues more often than the well-mixed reference. We call this phenomenon the *double-edged sword*: the same spatial structure provided by a biofilm reduces the rescue probability at low dose but raises it at high dose, and two conditions, an active interior and a high antibiotic dose, have to coincide before the biofilm rescues more often than the well-mixed reference at all.




Figure 2: Rescue probability for the biofilm (dark) versus the well-mixed reference (light), at low dose ($d_e = 1.05$, left columns) and high dose ($d_e = 2.8$, right columns). Panel A: dormant core ($b_c = 0$). Panel B: active core ($b_c = 0.8$). Each subplot uses its own y-axis scale. *Parameters:* $N_0 = 1000$, $l = 100$, $\mu = 10^{-4}$; each bar averages over 5,000 independent simulation runs.

The most pressing follow-up is what the active core is contributing at high dose that it fails to contribute at low dose, and why the dormant biofilm gets no benefit from its extra lifetime in the first place. Both questions concern where rescue actually comes from inside the biofilm. The next section opens up rescue itself by separating it into the two things any population needs in order to rescue: a resistance mutation has to arise (supply), and once it does, it has to grow to dominance in the drug-exposed edge, where antibiotic selection makes resistance an advantage (establishment). Tracking these two ingredients across the same conditions as Figure 2 is what lets us see where the biofilm and the well-mixed reference actually diverge.

4 Supply and establishment

Section 3 ended by separating rescue into two requirements: a resistance mutation has to arise (supply) before extinction, and once it does, it has to grow to dominance in the drug-exposed edge, where antibiotic selection makes resistance an advantage (establishment). To understand

where the dose-dependent flip in Figure 2 comes from, we measure each requirement separately. The comparison between biofilm and well-mixed then splits into two pieces: how many mutations each population produces, and what fraction of those mutations establish in the edge.

Before splitting rescue into its two pieces, we first check that the bar-plot flip is a continuous feature of the dose response, not an artifact of sampling at two specific doses. Figure 3 Panel A shows the three populations across the full dose sweep. The active-core biofilm here is drawn at our default $b_c = 0.2$, which sets the core's birth and death rate to one-fifth of the edge birth rate $b_e = 1$; this is considerably less active than the $b_c = 0.8$ used for Figure 2, where the core divides and dies at four-fifths of the edge rate. The active-core biofilm crosses the well-mixed reference smoothly near $d_e \approx 1.4$ and stays above it at every higher dose; the dormant-core biofilm ($b_c = 0$, no metabolism in the interior) sits below the well-mixed reference at every dose. Panel B plots the same data as differences from the well-mixed reference: the active biofilm's $\Delta P(\text{rescue})$ crosses zero near $d_e \approx 1.4$, and the dormant biofilm's $\Delta P(\text{rescue})$ stays below zero throughout the sweep. Both bar-plot patterns therefore persist across the full dose range, and the crossover is robust to the level of core activity: even at the more modest $b_c = 0.2$, an active interior is enough to flip the comparison at high dose.

Figure 4 Panel A reports the first piece of the decomposition: the difference in mutation count $\Delta S = S_{\text{BF}} - S_{\text{WM}}$ between the biofilm and the well-mixed reference, where S is the mean number of R mutations a population produces per simulation run. Two observations. First, the dormant-core biofilm sits essentially on the $y = 0$ reference line at every dose: the dormant biofilm and the well-mixed reference produce essentially the same number of mutations, even though the dormant biofilm lives longer. We return to why that extra lifetime does not translate into extra mutations in a later section. Second, the active-core biofilm sits above the $y = 0$ line at every dose: an actively dividing interior produces more mutations than the well-mixed reference at every dose. The number of extra mutations the active biofilm produces over the well-mixed reference shrinks with dose, from roughly 1.5 extra mutations per run at $d_e = 1.05$ down to ~ 0.03 at $d_e = 2.8$, because the absolute number of mutations any population produces shrinks at higher dose (shorter lifetimes give fewer cell divisions during which a mutation can occur). In ratio terms, however, the active biofilm produces about $1.7\times$ as many mutations as the well-mixed reference at every dose. Notably, the number of extra mutations the active biofilm produces does not track where the biofilm wins: the largest excess in mutation count is at low dose, precisely where the active biofilm rescues less often than the well-mixed reference (Figure 3 Panel B); the excess shrinks toward zero at high dose, where the biofilm actually rescues more often. Where the active biofilm has the most additional mutations available, it converts them into the fewest rescue events.

Figure 4 Panel B reports the second piece: the difference in per-mutation establishment between biofilm and well-mixed, $\Delta P(\text{est}) = P_{\text{est, BF}} - P_{\text{est, WM}}$, where $P(\text{est})$ is the fraction of edge-born R lineages whose populations reach the establishment threshold $k_{\text{est}} = 3$ cells before going extinct. Both biofilm curves sit below the $y = 0$ reference line at every dose: a given edge-born R lineage is consistently less likely to grow past the early-extinction bottleneck inside a biofilm than in the well-mixed reference, by roughly 0.05 to 0.10 in absolute probability across most of the dose range. This reduced establishment applies to both dormant and active biofilms and persists across the entire sweep. We return to why edge-born R lineages are less likely to escape early extinction inside a biofilm in a later section.




Figure 3: Probability of rescue versus dose for the three populations: well-mixed reference (open circles), biofilm with a dormant core ($b_c = 0$, open squares), and biofilm with an active core ($b_c = 0.2$, filled triangles). Panel A: absolute $P(\text{rescue})$, log y -axis. Panel B: difference relative to the well-mixed reference, $\Delta P(\text{rescue}) = P_{\text{BF}} - P_{\text{WM}}$, with the dotted horizontal line at zero marking the no-difference baseline. Each point averages 5,000 independent simulations. *Parameters:* $N_0 = 1000$, $l = 100$, $\mu = 10^{-4}$.




Figure 4: Decomposing the biofilm-to-well-mixed comparison into the two ingredients rescue requires. Panel A: difference in mutation supply $\Delta S = S_{\text{BF}} - S_{\text{WM}}$, where S is the mean number of R mutations a population produces over a simulation run. Symlog y -axis, with the dotted line at $y = 0$ marking equal supply. Panel B: difference in per-mutation establishment fraction for edge-born mutations, $\Delta P(\text{est}) = P_{\text{est, BF}} - P_{\text{est, WM}}$, where $P(\text{est})$ is the fraction of edge-born R lineages whose populations reach the establishment threshold $k_{\text{est}} = 3$ cells before going extinct. Linear y -axis, with the dotted line at $y = 0$ marking equal establishment. Both panels show two biofilm conditions across the dose sweep: dormant core ($b_c = 0$, open squares) and active core ($b_c = 0.2$, filled triangles). Each point averages 5,000 independent runs. *Parameters:* $N_0 = 1000$, $l = 100$, $\mu = 10^{-4}$.